

Statistical Inference Course Project - Part 1

Alex McBride

Friday, February 06, 2015

Contents

Synopsis	1
Create the Distribution	1

Synopsis

Investigate the exponential distribution in R and compare it with the Central Limit Theorem. The exponential distribution can be simulated in R with `rexp(n, lambda)` where `lambda` is the rate parameter. The mean of exponential distribution is $1/\lambda$ and the standard deviation is also $1/\lambda$. Set $\lambda = 0.2$ for all of the simulations. Investigate the distribution of averages of 40 exponentials running a thousand simulations.

Illustrate via simulation and associated explanatory text the properties of the distribution of the mean of 40 exponentials. You should 1. Show the sample mean and compare it to the theoretical mean of the distribution. 2. Show how variable the sample is (via variance) and compare it to the theoretical variance of the distribution. 3. Show that the distribution is approximately normal.

Create the Distribution

```
# set seed for reproducibility
set.seed(10)

# set lambda to 0.2
lambda <- 0.2

# 40 samples
n <- 40

# 1000 simulations
sim <- 1000

# simulate
sim_exp <- replicate(sim, rexp(n, lambda))

# calculate mean of exponentials
means_exp <- apply(sim_exp, 2, mean)
```

Show and compare the distribution mean and variance

Theoretical mean is $1/\lambda = 5$

Our analytical mean is `mean(means_exp) = 5.0450596`


The theoretical standard deviation is $1/\lambda/\sqrt{n} = 0.7905694$

Our analytical standard deviation is $\text{sd}(\text{means_exp}) = 0.7982821$

The theoretical variance is $((1/\lambda) * (1/\sqrt{n}))^2 = 0.625$

Our analytical variance is $\text{var}(\text{means_exp}) = 0.6372544$

```
hist(means_exp, xlab = "mean", main = "Exponential Function Simulations")
abline(v = mean(means_exp), col = "green")
abline(v = 1/lambda, col = "yellow")
```



Show that the distribution is approximately normal

```
expscale <- scale(means_exp)
hist(expscale, probability=T, main="Distribution Density", col = "cyan", ylim=c(0, 0.5))
lines(density(expscale), col = "purple", pch = 22, lty=5)
# Compare with the standard normal distribution
curve(dnorm(x,0,1), -4, 4, col="red", add=T)
```

